

Consumer-Grade AR Smart Glasses

The Evolution of Consumer-Grade AR Smart Glasses

Jin Wang, Seunggil Jeon, Gwanggil Jeon
Incheon National University

Abstract—The consumer AR landscape is shifting from passive screen-mirroring to autonomous intelligent assistants. This paper explores this evolution, highlighting the integration of spatial computing and Edge-AI within consumer electronics ecosystems. To address the challenge of high-performance AI on constrained hardware, we evaluate two deployment strategies: a native C++ approach (NCNN/YOLO26) for real-time object detection and Unity 6 for streamlined spatial integration. We further examine how vision-language foundation models are enabling richer multimodal scene understanding on edge devices. Additionally, we analyze agentic AI frameworks, exemplified by OpenClaw, that leverage LLMs to orchestrate autonomous, context-aware interactions. We conclude by outlining a roadmap for next-generation AR, where on-device intelligence and spatial awareness redefine the human-computer interface for global consumers.

■ **CONSUMER-GRADE** Augmented Reality (AR) glasses [1] are shifting from tethered display peripherals to intelligent, context-aware assistants. After a first decade dominated by optical efficiency and form-factor trade-offs (e.g., Birdbath and waveguide optics) [2], the focus has turned to Edge-AI integration, where Artificial Intelligence is now reshaping interactive AR experiences [3]. The central challenge is to execute complex computer-vision and generative-AI workloads within the strict thermal, privacy, and power envelopes of wearable form factors [4], [5].

Object detection remains the cornerstone of AR spatial awareness. Although the YOLO family [6], [7] is the de facto standard for real-time detection, edge

deployment requires a careful balance between speed and accuracy [8]. Recent implementations such as the DeepReality framework [9] and AI-powered first-aid systems [10] confirm that lightweight YOLO variants can sustain real-time inference on wearable platforms. In this paper we adopt *YOLO26* which is the latest publicly released iteration of the Ultralytics YOLO line (preview releases described in [11]–[13] during 2025–2026). It replaces the conventional Non-Maximum Suppression (NMS) step with an end-to-end, NMS-free decoding head and reports up to a 30% reduction in inference latency on mobile NPU/GPU hardware compared with its predecessors [12]. Combined with high-performance native kernels such as NCNN, this architecture enables a new class of “offline-first” AR deployments that bypass heavy cloud dependencies [14].

Digital Object Identifier 10.1109/MCE.YYYY.Doi Number

Date of publication DD MM YYYY; date of current version DD MM YYYY

The state-of-the-art is moving beyond reactive detection toward proactive assistance, driven by the convergence of XR and Large Language Models (LLMs) [15]. This integration has birthed Agentic AI in AR. Literature explores applications ranging from ubiquitous AGI companions [16] to multimodal assistants for visually impaired navigation [17], [18]. Meanwhile, as AR systems increasingly rely on visual perception for autonomous decision-making, ensuring robustness against visual information manipulation has emerged as a critical challenge [27].

Crucially, AR interfaces are transitioning from static displays to adaptive, sensible environments. Researchers are exploring unobtrusive interactions with proactive agents and memory-augmented models for deep spatiotemporal reasoning [25], [26]. To address privacy and latency, modern systems are pushing completely local voice-and-vision inference onto wearable nodes [19], though challenges remain in visual token compression [20], [21].

To bridge the gap between perception and inference without overloading edge devices, our architecture integrates an underlying autonomous agent logic layer. Multi-agent AI frameworks have already shown profound potential in adaptive AR [24]. Building upon this, we leverage the open-sourced OpenClaw framework [22] instead of designing a reasoning engine from scratch. Originally a local-first autonomous agent, Autonomous AI agent systems enable the AR pipeline to interpret rich semantic output from the YOLO26 multi-task network and orchestrate complex task workflows via natural language instructions [23].

In summary, this paper provides a magazine-style overview of the architectural evolution required for on-device AR intelligence and clarifies the design space for practitioners. **The key takeaways and contributions of this article are:**

- *Three-era roadmap.* We organize the evolution of consumer-grade AR glasses into three stages: display-centric, spatially aware, and cognitively intelligent. We identify the technology drivers behind each transition.
- *Closed-loop edge-AI architecture.* We propose and describe a heterogeneous three-layer pipeline (Perception–Processing–Interaction) that decouples the lightweight optical glass from a dedicated AI-compute unit.
- *Dual-track deployment study.* We articulate the trade-offs between a Native C++/NCNN track and

a Unity 6 / C# track. The Native C++/NCNN track is optimized for raw computational throughput and minimal inference latency and the Unity 6 / C# track is optimized for development velocity, mature XR rendering, and rapid spatial-UI integration.

- *Agentic AR illustration.* Through a concrete dietary-assistant use case built on YOLO26 detection and the OpenClaw agent framework, we show how AR glasses transition from passive viewers to proactive assistants that retrieve and present multimodal contextual knowledge.

THE EVOLUTION OF CONSUMER-GRADE SMART DEVICES

The trajectory of consumer-grade smart devices over the past two decades is not merely a story of increasing processing power, but a fundamental transformation in human-computer interaction (HCI) and environmental awareness. This evolution can be categorized into three distinct eras:

The Era of Mobile Connectivity: The inception of the modern smart-device era was marked by the smartphone revolution. During this decade, devices shifted from offline functional tools to always-connected portals; the primary focus was on app-based ecosystems and cloud synchronization [14]. However, the interaction remained largely two-dimensional, confined to glass touchscreens, and devices remained largely oblivious to the physical context of the user.

The Era of Pervasive Sensing and IoT: As silicon continued to drive miniaturization, intelligence began to leak out of the smartphone into a broader ecosystem of wearables and IoT devices. Smartwatches, hearables, and early-stage AR glasses introduced pervasive sensing, with on-board tracking of biometrics and inertial orientation (IMU) [4]. Yet, the intelligence was largely reactive, waiting for a voice command or a touch, and the AR component was primarily used for screen mirroring, acting as a secondary display rather than a spatial computer.

The Era of Spatial Intelligence and Edge-AI: We are currently witnessing the most significant leap—the transition to Spatial Intelligence. Driven by breakthroughs in Edge-AI [5], [14], smart devices are no longer just connected but truly *aware* of their surroundings.

- *From Pixels to Objects*—Modern devices, powered by high-performance inference engines such

as NCNN and Unity, can now perform real-time object detection (e.g., YOLO26) directly on-device [8], [11].

- *From Display to Assistant*—The integration of Large Language Models (LLMs) and Vision–Language Models (VLMs) has transformed devices from passive screens into proactive intelligent assistants [15], [21].
- *The Decoupling Trend*—A defining characteristic of this era is the split-processing architecture, where the display is kept lightweight for comfort while the computation is handled by a dedicated AI-compute unit [4].

Table 1 shows the summary table of the evolution of consumer-grade smart devices.

THE EVOLUTIONARY PATH OF AR WEARABLES: FROM PERIPHERALS TO COGNITIVE ENTITIES

The development of consumer-grade Augmented Reality (AR) glasses over the past decade can be viewed as a systematic progression toward balancing ergonomics, immersion, and intelligence. This evolution is characterized by three pivotal technological shifts.

The Display-Centric Stage

Early consumer AR devices were primarily designed as Head-Mounted Displays (HMDs). In this initial stage, the architecture was fundamentally a display-centric model: the device functioned as a tethered peripheral, using simple optical engines (such as early Birdbath or waveguide prisms [2]) to project a fixed-focus 2D virtual screen into the user's line of sight. These devices lacked environmental awareness; they were world-blind, with no on-board sensors to track spatial coordinates or recognize physical objects. The primary value proposition was portability—replacing a physical monitor with a wearable private screen.

The Spatial Awareness Stage

The second major evolution involved the integration of spatial computing. By incorporating localized sensors—including Inertial Measurement Units (IMUs) and dual-monochrome cameras—AR glasses transitioned from static displays to spatially-aware interfaces. This stage introduced Simultaneous Localization and Mapping (SLAM), enabling 6DoF (Six

Degrees of Freedom) tracking, which allowed virtual content to be anchored to physical surfaces and created the illusion of digital–physical coexistence [1], [2]. However, while the device understood *where* it was in space, it still did not understand *what* it was looking at: the intelligence remained purely geometric, not semantic.

The Cognitive Intelligence Stage

We are currently entering the third and most transformative stage: Cognitive AR, defined by the fusion of Edge-AI and agentic logic.

Semantic Perception: Through optimized neural-network deployment (using frameworks such as NCNN or Unity Sentsis), modern AR glasses can now perform real-time object detection and semantic segmentation [8], [9]. The device can identify objects, read text, and recognize human gestures.

Autonomous Interaction: The integration of Large Language Models (LLMs) and Vision–Language Models has shifted the paradigm from a command–response interaction to a proactive-assistant model [15], [16]. AR glasses are evolving into autonomous intelligent agents that can perceive the environment, reason about user intent, and provide context-aware information—e.g., real-time translation or step-by-step guidance—without explicit user triggers [23], [25].

PROPOSED SYSTEM ARCHITECTURE: A CLOSED-LOOP INFERENCE PIPELINE

To facilitate the transition from screen mirroring to intelligent assistance, we propose a heterogeneous edge-computing architecture tailored for consumer-grade split-processing AR glasses. As illustrated in Figure 1, the architecture establishes a seamless, low-latency closed-loop pipeline organized into three logical layers: Perception (Glass), Processing (Compute Unit), and Interaction (Unity rendering).

Rationale for the dual-track Processing Layer. The Processing Layer is intentionally split into two parallel tracks because no single runtime simultaneously optimizes both compute throughput and developer productivity on current AR hardware. The **Native C++ / NCNN track** is included to maximize raw computational performance. It gives direct control over INT8 quantization, NPU/GPU kernel dispatch, and memory layout, and minimizes end-to-end inference latency for time-critical detection. The **Unity 6 / C# track** is included to optimize development velocity and

TABLE 1. The Evolution of Consumer-Grade Smart Devices.

Feature	Era of Mobile Connectivity	Era of Pervasive Sensing & IoT	Era of Spatial Intelligence & Edge-AI
Core Devices	Smartphones	Smartwatches, Hearables, Early AR	Spatial Computers, AI-Compute Units
Awareness	Internet-Connected (Oblivious to physical context)	Basic Sensing (Biometrics, IMU orientation)	Spatial Awareness (Object & environment understanding)
Interaction	2D Glass Touchscreens	Reactive (Voice/Touch commands)	Proactive (Intelligent Assistants)
Key Tech	App Ecosystems, Cloud Sync	Miniaturization, Sensor Integration	Edge-AI, LLMs/VLMs, NCNN Inference
Visual Focus Architecture	2D Pixels Integrated (All-in-one)	Screen Mirroring Wearable + Smartphone Link	Spatial Computing Split-Processing (Decoupled)
Key Shift	From Offline to Online	From Mobile to Wearable	From “Pixels” to “Objects”

spatial-rendering quality. Unity 6 supplies a mature XR/SLAM toolchain (NRSDK), the Universal Render Pipeline (URP) for stereoscopic 3D output, and a high-level component model that accelerates iteration on spatial UI. The closed-loop pipeline therefore evaluates a concrete trade-of between peak inference performance (C++ track) versus integrated spatial rendering and faster iteration (Unity track). It makes the C++ and Unity interoperability boundary explicit through a JNI bridge.

Perception Layer (Sense)

The perception layer represents the input stage, leveraging the on-device sensors of the AR glass to capture raw environmental data. A wide-field RGB camera continuously streams video, acting as the primary visual input. This raw data stream is transmitted via a high-bandwidth, tethered connection (e.g., MIPI over USB-C) to the external compute unit, serving as the foundation for both cognitive understanding and spatial positioning.

Processing Layer (Think)

The compute unit (e.g., a high-performance computer unit or smartphone) serves as the brain of the system, employing the localized dual-track deployment strategy described above to maximize heterogeneous compute efficiency:

Native Track (C++): This track handles time-critical computer-vision tasks directly on mobile NPU/GPU hardware, where the C++/NCNN runtime is

selected specifically to maximize inference throughput and minimize per-frame latency [8], [14].

- *Input Pre-processing*—Raw images are resized and normalized to meet model requirements.
- *NCNN Inference Engine*—An optimized INT8-quantized YOLO26-nano model executes high-performance object detection, outputting rich semantic data (bounding boxes, classes, scores) for recognized physical objects [11]–[13].

Unity Track (C#): This track manages spatial awareness and model synchronization, where Unity 6 is selected specifically because its XR/SLAM toolchain and URP rendering accelerate development of spatial UIs.

- *XR NRSDK*—Processes sensor-fusion data to provide precise SLAM and 6DoF tracking information.
- *AI Model Manager*—Acts as a bridge, using a JNI/Native Plugin to pull the detection results (C++ structs) into the Unity C# environment and synchronize semantic data with current spatial coordinates.

Interaction Layer (Act)

The interaction layer, built upon the Unity 6 Universal Render Pipeline (URP), is responsible for the final spatial anchoring and stereoscopic visualization. The AR Spatial Anchoring Module is the critical component that fuses the asynchronous inputs: *Object ID & Coordinates* from the NCNN engine and *real-time SLAM data* from the NRSDK. It generates




FIGURE 1. The proposed heterogeneous edge-computing architecture and closed-loop inference pipeline for consumer-grade AR glasses. The Processing Layer is intentionally split into a Native C++ track (optimized for inference throughput) and a Unity 6 / C# track (optimized for XR rendering and development velocity).

persistent 3D virtual objects that are aligned with millimetre precision to the physical world, maintaining their coordinates over time.

The stereoscopic output is displayed on the AR glass, providing the user with an intuitive first-person AR view. Finally, a dedicated closed-loop feedback path continuously transmits refined sensor data and critical latency metrics back to the processing layer, allowing the system to dynamically optimize inference performance and spatial alignment to ensure a highly responsive and comfortable user experience.

Illustrative Use Case: A Proactive Dietary Assistant

The operational capabilities and the hierarchical progression of our proposed AR pipeline are visually demonstrated in Figure 2 and Figure 3.

To make the operation of the pipeline concrete, consider a single first-person scenario that exercises

every layer. A user wearing the AR glasses walks into a cafeteria and looks down at a buffet tray. (i) The Perception Layer streams the RGB feed of the tray over USB-C to the compute unit. (ii) In the Processing Layer, the Native C++ / NCNN track runs INT8 YOLO26-nano on each frame and outputs bounding boxes labeled *fish*, *radish*, *rice*, *soup*, and *tomato*, while the Unity 6 / C# track produces SLAM coordinates and synchronizes the detections with the 3D anchor pose via the JNI bridge. (iii) The OpenClaw agent layer interprets this semantic stream, queries a small on-device knowledge module for the calorie and macronutrient profile of each item, and assembles a short summary. (iv) The Interaction Layer renders the summary as a stereoscopic floating label anchored to the tray (Figure 3); when the user glances at a different dish, the closed-loop feedback path triggers re-anchoring and content refresh in real time. This single example shows how detection (Figure 2), spatial



FIGURE 2. Baseline real-time object detection from the first-person perspective (FPV) of the AR glasses.




FIGURE 3. Demonstration of agentic AR capabilities through context-aware food classification and proactive multimodal assistance.

anchoring, and agentic reasoning combine to turn the glasses from a passive viewer into a proactive dietary assistant.

Figure 2 illustrates the baseline real-time object detection performance from the first-person perspective (FPV) of the AR glasses. Powered by the optimized YOLO26 inference engine, the system accurately detects and localizes everyday environmental elements. Bounding boxes are dynamically rendered around persons, tables, and chairs within the user's field of view, validating the system's capability to maintain continuous, low-latency spatial awareness in a complex physical environment.

Building upon this foundational perception, Fig-

ure 3 exemplifies the system's advanced agentic capabilities in a practical scenario. Here, the pipeline goes beyond simple bounding-box localization to perform fine-grained classification of various food items. Crucially, it retrieves and renders detailed contextual information about the recognized food—specific categories, nutritional facts, and dietary profiles—displayed seamlessly in the top-left corner of the AR interface. This scenario demonstrates the paradigm shift from passive visual perception to proactive, multimodal intelligent assistance, providing the user with actionable insights directly within their line of sight.

CONCLUSION AND FUTURE WORK

Conclusion

This paper has explored the transformative journey of consumer-grade AR glasses, tracing their evolution from simple head-mounted displays to intelligent, spatially-aware assistants. Our research demonstrates that the transition from Screen Mirroring to Intelligent Assistance is primarily driven by the successful integration of high-performance Edge-AI within the constraints of mobile form factors.

Through experimental benchmarks on leading hardware platforms, we have validated that native deployment frameworks such as NCNN, combined with optimized neural networks such as YOLO26, can sustain real-time, low-latency inference on consumer-grade AR hardware. This performance is critical for maintaining the 3D spatial consistency required for meaningful AR interactions. Furthermore, the emergence of engine-integrated AI via Unity Sentis in the Unity 6 environment has significantly lowered the barrier for developers to create semantically-aware AR applications. Ultimately, our findings confirm that on-device cognitive intelligence is not only feasible but essential for the next generation of consumer electronics.

Future Work: Toward Agentic AR and Multi-modal Intelligence

While current systems excel at reactive object detection, the future of AR lies in proactive autonomy. Our future research will focus on several key frontiers, spanning core intelligence, user interaction, and the surrounding ecosystem.

Integration of Large Language Models (LLMs): By embedding lightweight, on-device LLMs, AR wearables can evolve into autonomous agents. Building

upon frameworks like OpenClaw, these agents will be capable of interpreting complex user intents and delivering multi-step, real-time guidance.

Multi-modal Interaction Fusion: Future iterations will investigate the synergy of computer vision, spatial audio processing, and fine-grained gesture recognition. This human-centric paradigm will foster more intuitive interfaces, empowering the assistant to proactively anticipate user needs through a holistic comprehension of the surrounding environment.

Collaborative Edge-Cloud Computing: Anticipating the maturation of 5G/6G communication networks, we will explore hybrid computational architectures that harmonize the privacy and low-latency advantages of Edge-AI with the expansive knowledge-retrieval capacities of cloud infrastructure, ensuring robust intelligence in information-dense scenarios.

Next-Generation User Interaction: Beyond touch and voice, AR glasses will increasingly rely on eye gaze, fine-grained hand and finger tracking, and electromyographic (EMG) wristbands as primary input channels. We expect future work to focus on attention-aware, gaze-driven interfaces in which the agent infers *what* the user is asking about from *what* they are looking at, reducing the cognitive load of explicit commands.

Ecosystem and Standardization Trends: Finally, scalable consumer AR depends on a healthy developer ecosystem. We anticipate three converging trends: (i) cross-vendor XR runtime standards (e.g., OpenXR-style abstractions) that let a single agent pipeline target multiple glass form factors; (ii) on-device model marketplaces and shared agent skills, so that developers can publish reusable detection and reasoning modules rather than retraining from scratch; and (iii) tighter integration with consumer-electronics platforms (smartphones, smartwatches, home IoT), with the AR glasses acting as a thin, spatially-aware front-end to an already pervasive ambient-computing stack.

By effectively bridging the chasm between raw perception and cognitive reasoning, the ensuing generation of AR smart glasses will transcend their current status as mere technological gadgets, solidifying their role as indispensable, ubiquitous companions in the daily lives of global consumers.

REFERENCES

1. R. Azuma, "A Survey of Augmented Reality," *Presence: Teleoperators & Virtual Environments*, vol. 6, no. 4, pp. 355–385, 1997.
2. J. Xiong, E.-L. Hsiang, Z. He, T. Zhan, and S.-T. Wu, "Augmented Reality and Virtual Reality Displays: Emerging Technologies and Future Perspectives," *Light: Science & Applications*, vol. 10, Art. no. 216, Oct. 2021, doi: 10.1038/s41377-021-00658-8.
3. L. Kugler, "How AI Could Supercharge AR and VR," *Communications of the ACM*, vol. 68, no. 10, pp. 15–17, Oct. 2025, doi: 10.1145/3748667.
4. X. Wang, Z. Tang, J. Guo, T. Meng, C. Wang, T. Wang, and W. Jia, "Empowering Edge Intelligence: A Comprehensive Survey on On-Device AI Models," *ACM Computing Surveys*, vol. 57, no. 9, art. 228, pp. 1–39, Apr. 2025, doi: 10.1145/3724420.
5. V. J. Reddi, "Generative AI at the Edge: Challenges and Opportunities: The next phase in AI deployment," *ACM Queue*, vol. 23, no. 2, pp. 79–137, May 2025, doi: 10.1145/3733702.
6. C.-Y. Wang, A. Bochkovskiy, and H.-Y. M. Liao, "YOLOv7: Trainable Bag-of-Freebies Sets New State-of-the-Art for Real-Time Object Detectors," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, Vancouver, BC, Canada, 2023, pp. 7464–7475, doi: 10.1109/CVPR52729.2023.00721.
7. Y. Feng, J. Huang, S. Du, S. Ying, J.-H. Yong, Y. Li, G. Ding, R. Ji, and Y. Gao, "Hyper-YOLO: When Visual Object Detection Meets Hypergraph Computation," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 47, no. 4, pp. 2388–2401, Apr. 2025, doi: 10.1109/TPAMI.2024.3524377.
8. P. Mittal, "A comprehensive survey of deep learning-based lightweight object detection models for edge devices," *Artif. Intell. Rev.*, vol. 57, no. 9, Sep. 2024, Art. no. 242, doi: 10.1007/s10462-024-10877-1.
9. R. Pierdicca, F. Tonetto, M. Paolanti, M. Mameli, R. Rosati, and P. Zingaretti, "DeepReality: An open source framework to develop AI-based augmented reality applications," *Expert Syst. Appl.*, vol. 249, Sep. 2024, Art. no. 123530, doi: 10.1016/j.eswa.2024.123530.
10. M. Abo-Zahhad, M. N. Zakaria, F. M. Sharaf, M. M. Ismaiel, H. Hafrag, and Y. M. Amer, "Development of an AI-powered AR glasses system for real-time first aid guidance in emergency situations," *BioData Min.*, vol. 18, Art. no. 59, Aug. 2025, doi: 10.1186/s13040-025-00473-6.
11. P. Hidayatullah and R. Tubagus, "YOLO26: A Comprehensive Architecture Overview and Key Improvements," arXiv preprint arXiv:2602.14582, Feb. 2026.
12. R. Sapkota, R. H. Cheppally, A. Sharda, and M.

- Karkee, "YOLO26: Key Architectural Enhancements and Performance Benchmarking for Real-Time Object Detection," arXiv preprint arXiv:2509.25164, 2025.
13. S. Chakrabarty, "YOLO26: An Analysis of NMS-Free End to End Framework for Real-Time Object Detection," arXiv preprint arXiv:2601.12882, Jan. 2026.
14. J.-H. Syu, J. C.-W. Lin, G. Srivastava, and K. Yu, "A Comprehensive Survey on Artificial Intelligence Empowered Edge Computing on Consumer Electronics," *IEEE Trans. Consum. Electron.*, vol. 69, no. 4, pp. 1023–1034, Nov. 2023, doi: 10.1109/TCE.2023.3318150.
15. J. Wang, W. Meng, Y. Jin, J. Wang, Y. Qin, and M. Yang, "A Survey of Visual-Language Foundation Models for Enhancing Virtual and Augmented Reality Interactivity," *Vis. Comput.*, vol. 42, Art. no. 127, Jan. 2026, doi: 10.1007/s00371-025-04332-7.
16. D. Zhang, Y. Li, Z. He, and X. Li, "Empowering smart glasses with large language models: Towards ubiquitous AGI," in *Adjunct Proc. 2024 ACM Int. Joint Conf. Pervasive Ubiquitous Comput. (UbiComp)*, Oct. 2024, pp. 631–633, doi: 10.1145/3675094.3678992.
17. Z.-Y. Lee, Y.-C. Sun, Y.-N. Wong, and C.-H. Hsu, "Exploring LLM-based assistants with smart glasses for the visually impaired," in *Proc. 10th ACM/IEEE Symp. Edge Comput. (SEC)*, Dec. 2025, Art. no. 62, pp. 1–8, doi: 10.1145/3769102.3774892.
18. A. Singh, M. A. Bhanushali, J. Luo, G. Luo, and S. Pundlik, "Assisting the blind to reach daily objects using smart glasses," *Displays*, vol. 93, Jul. 2026, Art. no. 103390, doi: 10.1016/j.displa.2026.103390.
19. Y. Tussa, A. Heredia, and N. Roy, "Lessons Learned from Developing a Privacy-Preserving Multimodal Wearable for Local Voice-and-Vision Inference," in *Proc. ACM Conf. Embedded Networked Sensor Syst. (SenSys '26)*, Saint-Malo, France, May 2026.
20. P. Xu, X. Zhu, and D. A. Clifton, "Multimodal Learning with Transformers: A Survey," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 45, no. 10, pp. 12113–12132, Oct. 2023, doi: 10.1109/TPAMI.2023.3275156.
21. J. Zhang, J. Huang, S. Jin, and S. Lu, "Vision-Language Models for Vision Tasks: A Survey," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 46, no. 8, pp. 5625–5644, Aug. 2024, doi: 10.1109/TPAMI.2024.3369699.
22. P. Steinberger, "OpenClaw: An Autonomous Open-Source AI Agent," GitHub repository, Jan. 2026. [Online]. Available: <https://github.com/openclaw/openclaw>
23. D. B. Acharya, K. Kuppan, and B. Divya, "Agentic AI: Autonomous Intelligence for Complex Goals—A Comprehensive Survey," *IEEE Access*, vol. 13, pp. 18912–18936, 2025, doi: 10.1109/ACCESS.2025.3532853.
24. N. Leins, J. Gonnermann-Müller, M. Teichmann, and S. Pokutta, "Beyond Static Instruction: A Multi-agent AI Framework for Adaptive Augmented Reality Robot Training," in *Companion Proc. 21st ACM/IEEE Int. Conf. Hum.-Robot Interact. (HRI '26)*, 2026, pp. 989–993, doi: 10.1145/3776734.3794542.
25. G. Lee, M. Xia, N. Numan, X. Qian, D. Li, Y. Chen, A. Kulshrestha, I. Chatterjee, Y. Zhang, D. Manocha, D. Kim, and R. Du, "Sensible Agent: A Framework for Unobtrusive Interaction with Proactive AR Agents," in *Proc. 38th Annu. ACM Symp. User Interface Softw. Technol. (UIST '25)*, Busan, Republic of Korea, Sept. 2025, doi: 10.1145/3746059.3747748.
26. D. Choi, T. Kwon, D. Yang, H. Kim, and J. Yeo, "Designing Memory-Augmented AR Agents for Spatiotemporal Reasoning in Personalized Task Assistance," in *Proc. 2025 IEEE Int. Symp. Mixed Augmented Reality Adjunct (ISMAR-Adjunct)*, Oct. 2025.
27. Y. Xiu and M. Gorlatova, "Detecting Visual Information Manipulation Attacks in Augmented Reality: A Multimodal Semantic Reasoning Approach," *IEEE Trans. Vis. Comput. Graphics*, 2025. (Presented at IEEE ISMAR 2025; Best Paper Award.)

Jin Wang is with Incheon National University, Incheon, South Korea. Jin Wang received a BS in mathematics and applied mathematics from Zhejiang University in 2007, and an MS and PhD in electronic communications engineering from Hanyang University in 2009 and 2013. She was a professor at Xidian University. She is a professor at Incheon National University.

Gwanggil Jeon is with Incheon National University, Incheon, South Korea. Dr. Jeon received the B.S., M.S., and Ph.D. (summa cum laude) degrees from the Department of Electronics and Computer Engineering, Hanyang University, Seoul, Korea, in 2003, 2005, and 2008, respectively. From 2009.09 to 2011.08, he was with the School of Information Technology and Engineering, University of Ottawa.

Seunggil Jeon received the B.S., M.S., and Ph.D. degrees from the Department of Electronics and Computer Engineering, Hanyang University, Seoul, Korea, in 2002, 2004, and 2008, respectively. From 2009 to 2024, he was with the Samsung Electronics.